

# Sequence Labelling and Part-of-Speech Tagging

## COM6513 Natural Language Processing

Nikos Aletras

`n.aletras@sheffield.ac.uk`

`@nikalettras`

Computer Science Department

Week 4

Spring 2021



In the previous lecture...

- Our first **sequence modelling** problem: **Language Modelling**

# In this lecture...

- What about if we want to **assign a label to each word in a sequence?**

# In this lecture...

- What about if we want to **assign a label to each word in a sequence?**
- Sequence labelling!

# In this lecture...

- What about if we want to **assign a label to each word in a sequence?**
- Sequence labelling!
- Applications?

# Applications

- Part-of-Speech (POS) Tagging

$(\mathbf{x}, \mathbf{y}) = ([I, \textit{studied}, \textit{in}, \textit{Sheffield}],$   
 $[\textit{Pronoun}, \textit{Verb}, \textit{Preposition}, \textit{ProperNoun}])$

# Applications

- Part-of-Speech (POS) Tagging

$$(\mathbf{x}, \mathbf{y}) = ([I, \textit{studied}, \textit{in}, \textit{Sheffield}], \\ [\textit{Pronoun}, \textit{Verb}, \textit{Preposition}, \textit{ProperNoun}])$$

- Named Entity Recognition

$$(\mathbf{x}, \mathbf{y}) = ([\textit{Giannis}, \textit{Antetokounmpo}, \textit{plays}, \textit{for}, \textit{the}, \textit{Bucks}], \\ [\textit{Person}, \textit{Person}, \textit{NotEnt}, \textit{NotEnt}, \textit{NotEnt}, \textit{Org}])$$

- Machine Translation (reconstruct word alignments)

$$(\mathbf{x}, \mathbf{y}) = ([\textit{la}, \textit{maison}, \textit{bleu}], \\ [\textit{the}, \textit{house}, \textit{blue}])$$

We will use POS tagging as a running example

# Parts of Speech (POS)

Label words according to their syntactic function in a sentence:

The	results	appear	in	today	's	news
determiner	noun	verb	preposition	noun	possessive	noun



# Parts of Speech (POS)

Label words according to their syntactic function in a sentence:

The	results	appear	in	today	's	news
determiner	noun	verb	preposition	noun	possessive	noun

What could they be useful for?

# Parts of Speech (POS)

Label words according to their syntactic function in a sentence:

The	results	appear	in	today	's	news
determiner	noun	verb	preposition	noun	possessive	noun

What could they be useful for?

- text classification
- language modelling
- syntactic parsing
- named entity recognition
- question answering

# PoS Tags

- **Open** class:  
nouns, verbs, adjectives
- **Closed** class:  
determiners, prepositions, conjunctions, etc

# PoS definitions


- Most research uses the Penn Treebank PoS tag set
- Includes 45 tags making distinctions between:
  - verbs in active vs past tense
  - nouns in singular vs plural number
  - etc.

# PoS definitions

- Most research uses the Penn Treebank PoS tag set
- Includes 45 tags making distinctions between:
  - verbs in active vs past tense
  - nouns in singular vs plural number
  - etc.
- Penn Tree Bank inspired by English. Recent work has focused on the Universal PoS tag set:
  - 17 coarse tags: one noun class, one verb class, etc.
  - developed considering 22 languages

# Sequence labelling: Problem Setup

Data consists of word sequences with label sequences:


$$D_{train} = \{(\mathbf{x}^1, \mathbf{y}^1) \dots (\mathbf{x}^M, \mathbf{y}^M)\}$$

$$\mathbf{x}^m = [x_1, \dots, x_N]$$

$$\mathbf{y}^m = [y_1, \dots, y_N]$$

Learn a model  $f$  that predicts the best label sequence:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} f(\mathbf{x}, \mathbf{y}) \quad (1)$$

$\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}$  is the set of all possible combinations of label sequences  
and  $\mathcal{Y} = \{A, B, C \dots\}$  are the possible classes for each word.

# Could we use a dictionary-based model?

$\{the : determiner, can : modal, fly : verb\}$

# Could we use a dictionary-based model?

$\{the : determiner, can : modal, fly : verb\}$

Yes, but the same word can have different tags in different contexts.

I	can	fly
pronoun	modal	verb

vs:

I	can	fly
pronoun	verb	noun

**can** and 11.5% of the words in the Brown corpus have more than one tag



# Can we use a Markov model?

Use tags  $\mathbf{y}$  instead of words:

$$P(\mathbf{y}) = \prod_{n=1}^N P(y_n | y_{n-1})$$

# Can we use a Markov model?

Use tags  $\mathbf{y}$  instead of words:

$$P(\mathbf{y}) = \prod_{n=1}^N P(y_n | y_{n-1})$$

Our training data should be able to tell us that tagging is unlikely:

I	can	fly
pronoun	modal	noun

# Can we use a Markov model?

Use tags  $\mathbf{y}$  instead of words:

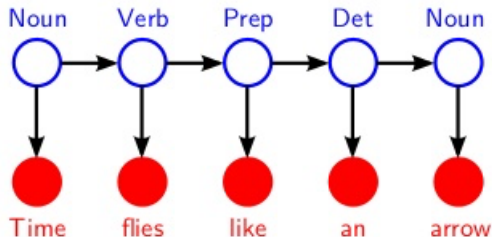
$$P(\mathbf{y}) = \prod_{n=1}^N P(y_n | y_{n-1})$$

Our training data should be able to tell us that tagging is unlikely:

I	can	fly
pronoun	modal	noun

What about the words? We will get the same  $N$ -tag long sequence for any sentence!

# Hidden Markov Model (HMM)



- Labels  $y_i$  (i.e. PoS tags) are hidden states emitting words.
- Assumptions:
  - 1st order Markov among the POS tags (current tag depends only on previous tag) *verb* *noun*
  - Each word only depends on its POS tag

# HMM: Derivation

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{y}|\mathbf{x}) \quad (\text{Bayes rule})$$

# HMM: Derivation

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{y}|\mathbf{x}) \quad (\text{Bayes rule})$$

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \frac{P(\mathbf{x}|\mathbf{y})P(\mathbf{y})}{\cancel{P(\mathbf{x})}} \quad (\text{word probabilities are constant})$$

# HMM: Derivation

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{y}|\mathbf{x}) \quad (\text{Bayes rule})$$

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \frac{P(\mathbf{x}|\mathbf{y})P(\mathbf{y})}{\cancel{P(\mathbf{x})}} \quad (\text{word probabilities are constant})$$

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{x}|\mathbf{y})P(\mathbf{y}) \quad (\text{1st order Markov})$$

# HMM: Derivation

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{y}|\mathbf{x}) \quad (\text{Bayes rule})$$

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \frac{P(\mathbf{x}|\mathbf{y})P(\mathbf{y})}{\cancel{P(\mathbf{x})}} \quad (\text{word probabilities are constant})$$

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P(\mathbf{x}|\mathbf{y})P(\mathbf{y}) \quad (\text{1st order Markov})$$

$$\hat{\mathbf{y}} \approx \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \prod_{n=1}^N P(x_n|y_n)P(y_n|y_{n-1})$$



# HMM: Training

- Maximum likelihood estimation (i.e. counts!):

$$P(y_n|y_{n-1}) = \frac{c(y_n, y_{n-1})}{c(y_{n-1})} \quad (\text{transition probabilities})$$

$$P(x_n|y_n) = \frac{c(x_n, y_n)}{c(y_n)} \quad (\text{emission probabilities})$$

# HMM: Training

- Maximum likelihood estimation (i.e. counts!):

$$P(y_n|y_{n-1}) = \frac{c(y_n, y_{n-1})}{c(y_{n-1})} \quad (\text{transition probabilities})$$

$$P(x_n|y_n) = \frac{c(x_n, y_n)}{c(y_n)} \quad (\text{emission probabilities})$$

- We can easily compute counts  $c(\cdot)$  using a labelled corpus (pairs of words-POS tags).

# HMM: Example

$\mathbf{x} = [\text{START}, I, \text{can}, \text{fly}, \text{END}]$

$\mathbf{y} = [\text{START}, \text{PPSS}, \text{MD}, \text{NN}, \text{END}]$

$$\hat{\mathbf{y}} \approx \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \prod_{n=1}^N P(x_n | y_n) P(y_n | y_{n-1})$$

$$\begin{aligned} P(\mathbf{y} | \mathbf{x}) = & P(I | \text{PPSS}) P(\text{PPSS} | \text{START}) \\ & P(\text{can} | \text{MD}) P(\text{MD} | \text{PPSS}) \\ & P(\text{fly} | \text{NN}) P(\text{NN} | \text{MD}) \end{aligned}$$

# Decoding/Inference

- So we have everything we need to decode/infer the most likely tag sequence for a sentence:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \prod_{n=1}^N P(x_n | y_n) P(y_n | y_{n-1})$$

- Just enumerate all possible tag sequences?

# Decoding/Inference

- So we have everything we need to decode/infer the most likely tag sequence for a sentence:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \prod_{n=1}^N P(x_n | y_n) P(y_n | y_{n-1})$$

- Just enumerate all possible tag sequences?  
Intractable. We would need to evaluate  $|\mathcal{Y}|^{\mathcal{N}}$  sequences!

# Decoding/Inference

- So we have everything we need to decode/infer the most likely tag sequence for a sentence:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \prod_{n=1}^N P(x_n | y_n) P(y_n | y_{n-1})$$

- Just enumerate all possible tag sequences?  
**Intractable.** We would need to evaluate  $|\mathcal{Y}|^{\mathcal{N}}$  sequences!
- We will see later how to decode efficiently!

# HMMs: Some extra points

- Higher order HMMs:
  - longer contexts, more expensive inference
  - benefits are usually small

# HMMs: Some extra points

- Higher order HMMs:
  - longer contexts, more expensive inference
  - benefits are usually small
- Smoothing:
  - what happens when we have unseen word/tags or tag-tag combinations?



# HMMs: Some extra points

- Higher order HMMs:
  - longer contexts, more expensive inference
  - benefits are usually small
- Smoothing:
  - what happens when we have unseen word/tags or tag-tag combinations?  
Use methods we learned in the language modeling lecture!

$$\begin{array}{r} C + k \\ \hline C + 1 \times k \end{array}$$

# HMMs: Limitations

- They generate probabilities for words and labels, we just want labels
- No overlapping features (e.g. unigrams+bigrams)
- No subword features (e.g. suffixes)

# Conditional Random Fields: Extend LR for Sequence Labelling

- Logistic regression can also provide probabilities but supports more flexible representations!

# Conditional Random Fields: Extend LR for Sequence Labelling

- Logistic regression can also provide probabilities but supports more flexible representations!
- Given a word, a candidate label of the current and the label of the previous word, predict the most likely label for that word using a (multi-class LR) in each time step → Conditional Random Fields

# Conditional Random Fields: Extend LR for Sequence Labelling

- Logistic regression can also provide probabilities but supports more flexible representations!
- Given a word, a candidate label of the current and the label of the previous word, predict the most likely label for that word using a (multi-class LR) in each time step → Conditional Random Fields

# Conditional Random Fields: Extend LR for Sequence Labelling

- Logistic regression can also provide probabilities but supports more flexible representations!
- **Given** a word, a candidate label of the current and the label of the previous word, predict the most likely label for that word using a (multi-class LR) in each time step → Conditional Random Fields
- CRF paper more than 11K citations since 2001, 10 year test of time award at ICML conference

# Conditional Random Fields

Decompose the per sentence  $\mathbf{x} = [x_1, \dots, x_N]$  prediction:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^N} f(\mathbf{x}, \mathbf{y})$$

into each word  $x_n$ :

$$\hat{y}_n = \arg \max_{y \in \mathcal{Y}} f(x_n; y_{n-1}, n) = \arg \max_{y \in \mathcal{Y}} \mathbf{w}^y \phi(x_n, y_{n-1}, n)$$

# Conditional Random Fields

Decompose the per sentence  $\mathbf{x} = [x_1, \dots, x_N]$  prediction:

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y} \in \mathcal{Y}^N} f(\mathbf{x}, \mathbf{y})$$

into each word  $x_n$ :

$$\hat{y}_n = \arg \max_{y \in \mathcal{Y}} f(x_n; y_{n-1}, n) = \arg \max_{y \in \mathcal{Y}} \mathbf{w}^y \phi(x_n, y_{n-1}, n)$$

- How to construct a feature vector  $\phi(x_n, y_n, y_{n-1}, n)$ ?



# CRF: Feature Vectors

- $\phi_1(x_n, y_n, y_{n-1}, n) = 1$  if  $y_n = ADVERB$  and the  $n$ -th word ends in “-ly”; 0 otherwise.  
“usually”, “casually”
- $\phi_2(x_n, y_n, y_{n-1}, n) = 1$  if  $n = 1, y_n = VERB$ , and the sentence ends in a question mark; 0 otherwise.  
“Is it true?”
- etc.

# CRF: Inference

The normalisation factor has to score all possible label sequences for all sentences, so it is ignored:

$$\arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} P_{CRF}(\mathbf{y} | \mathbf{x}; \mathbf{w}) = \arg \max_{\mathbf{y} \in \mathcal{Y}^{\mathcal{N}}} \sum_{n=1}^N \mathbf{w} \cdot \phi(y_n, y_{n-1}, \mathbf{x}, n)$$

# CRF: Training

- Training by minimising the negative log-likelihood objective:

$$\mathbf{w} = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \sum_{m=1}^M -\log P_{CRF}(\mathbf{y}^m | \mathbf{x}^m; \mathbf{w})$$

- using **Stochastic Gradient Descent**

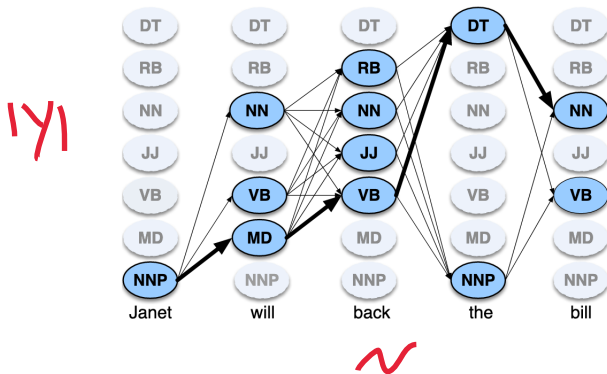
# Decoding with Viterbi

- Enumerating all possible tag sequences in HMM and CRF is intractable!

# Decoding with Viterbi

- Enumerating all possible tag sequences in HMM and CRF is intractable!
- Dynamic programming: store and re-use calculations
- Possible due to independence assumptions
- Keep track of the highest probability to reach each PoS tag for each word and how we got there

# Decoding with Viterbi



# Viterbi: Data structures

- **Viterbi score matrix**  $V^{|\mathcal{Y}| \times N}$ :

# Viterbi: Data structures

- **Viterbi score matrix**  $V^{|\mathcal{Y}| \times N}$ :
  - Tag set  $\mathcal{Y}$ , sentence  $\mathbf{x} = [x_1, \dots, x_N]$



# Viterbi: Data structures

- **Viterbi score matrix**  $V^{|\mathcal{Y}| \times N}$ :
  - Tag set  $\mathcal{Y}$ , sentence  $\mathbf{x} = [x_1, \dots, x_N]$
  - each cell contains the highest prob. for word  $n$  with tag  $y$

# Viterbi: Data structures

- **Viterbi score matrix**  $V^{|\mathcal{Y}| \times N}$ :

- Tag set  $\mathcal{Y}$ , sentence  $\mathbf{x} = [x_1, \dots, x_N]$
- each cell contains the highest prob. for word  $n$  with tag  $y$
- 1st order Markov: only depends on the previous tag  $y_{n-1}$   
$$V[y, n] = \max_{y_{n-1} \in \mathcal{Y}} V[y_{n-1}, n-1] \times P(y_n | x_n, y_{n-1})$$

# Viterbi: Data structures

- **Backpointer matrix**  $backptr^{|\mathcal{Y}| \times N}$ :

- **Backpointer matrix**  $backptr^{|\mathcal{Y}| \times N}$ :
  - instead of the max score, keep the previous tag that got it

# Viterbi: Data structures

- **Backpointer matrix**  $backptr^{|\mathcal{Y}| \times N}$ :
  - instead of the max score, keep the previous tag that got it
  - *argmax* instead of *max*  
 $backptr[y, n] = \arg \max_{y' \in \mathcal{Y}} V[y', n - 1] \times P(y|y') \times P(x_n|y)$

# Viterbi algorithm

**Input:** word sequence  $\mathbf{x} = [x_1, \dots, x_N]$ ,

$P(y_n|x_n, y_{n-1})$  probs

set matrix  $V^{|\mathcal{Y}| \times N} = 1$

**for**  $n = 1$  **to**  $N$  **do**

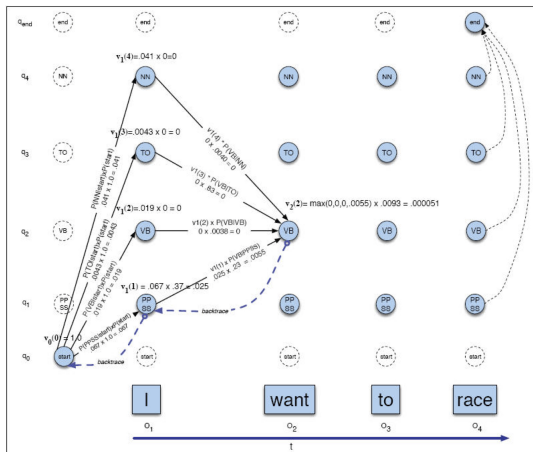
**for**  $y \in \mathcal{Y}$  **do**

$$V[y, n] = \max_{y_{n-1} \in \mathcal{Y}} V[y_{n-1}, n-1] \times P(y_n|x_n, y_{n-1})$$

$$\text{backptr}[y, n] = \arg \max_{y_{n-1} \in \mathcal{Y}} V[y_{n-1}, n-1] \times P(y_n|x_n, y_{n-1})$$

$$\text{backptr}[\text{None}, N+1] = \arg \max_{y_{n-1} \in \mathcal{Y}} V[y_{n-1}, N] \times P(\text{None}|y_{n-1})$$

# Viterbi diagram



Break the large arg max into smaller ones, left-to-right (**dynamic programming**)

# Beam Search: Inexact Decoding

- Viterbi performs exact search (under assumptions) by evaluating all options.



## Beam Search: Inexact Decoding

- Viterbi performs exact search (under assumptions) by evaluating all options.
- Get faster by being inexact, i.e. avoid labelling some candidate sequences with **Beam Search**

# Beam Search

- Do Viterbi, but keep only best  $k$  hypotheses at each step

# Beam Search

- Do Viterbi, but keep only best  $k$  hypotheses at each step
- If beam size is 1, then we have greedy search

# Beam Search

- Do Viterbi, but keep only best  $k$  hypotheses at each step
- If beam size is 1, then we have greedy search
- Often beams less than 10 get close to exact search, but much faster

# Beam Search

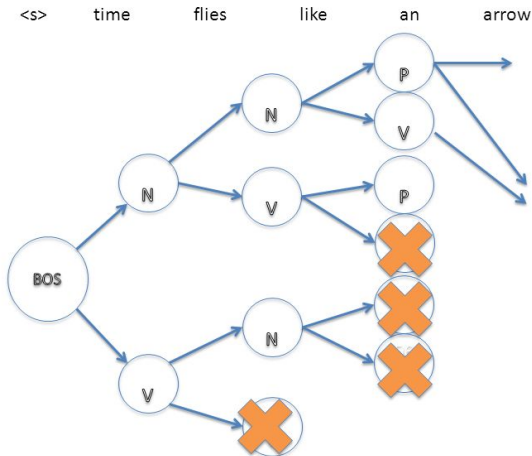
- Do Viterbi, but keep only best  $k$  hypotheses at each step
- If beam size is 1, then we have greedy search
- Often beams less than 10 get close to exact search, but much faster
- Beams must be of the same length to be comparable

# Beam Search

- Do Viterbi, but keep only best  $k$  hypotheses at each step
- If beam size is 1, then we have greedy search
- Often beams less than 10 get close to exact search, but much faster
- Beams must be of the same length to be comparable
- Attractive when we need complex feature functions i.e. avoid Markov assumptions)

# Beam Search: Example

## Beam Search, $k=3$



# Beam Search: Algorithm

**Input:** word sequence  $\mathbf{x} = [x_1, \dots, x_N]$ , weights  $\mathbf{w}$   
set beam  $B = \{(\mathbf{y}_{\text{temp}} = [\text{START}], \text{score} = 0)\}$ , size  $k$   
**for**  $n = 1$  **to**  $N$  **do**  
     $B' = \{\}$   
    **for**  $b \in B$  **do**  
        **for**  $y \in \mathcal{Y}$  **do**  
             $B' = B' \cup ([b.\mathbf{y}_{\text{temp}}; y], P([b.\mathbf{y}_{\text{temp}}; y] | x_n))$   
     $B = \text{TOP-}k(B')$   
**return**  $\text{TOP-}1(B)$



# Bibliography

- Chapter 8 from Jurafsky and Martin
- Sections 7.1-7.4, 7.5.3 and Chapter 8 from Eisenstein
- This [blog post on CRFs](#) by Edwin Chen
- Tutorial on CRFs by Sutton and McCallum

# Coming up next...

The best-studied, more complex than sequence labeling problem in NLP: **dependency parsing**